


# Time Series Analysis

## References:

- [1] F. Anderson, Time Series Analysis by Example : Hands on approach using R
- [2] M. O. Adenomon, An Introduction to Univariate and Multiple Time Series Analysis with R (Lambert Academic Pub, Saarbrücken, 2015)

## Identifying Explosions and Earthquakes Data recorded at Nilore, Pakistan



1. Introduction to Time Series
2. Creating Time Series Objects
3. Time Series Visualization
4. Time Series Stationarity
5. Time Series and Forecasting
6. Financial Time Series

# 1. Introduction to Time Series

- ▶ A **time series** is a sequence of data points ordered in time.

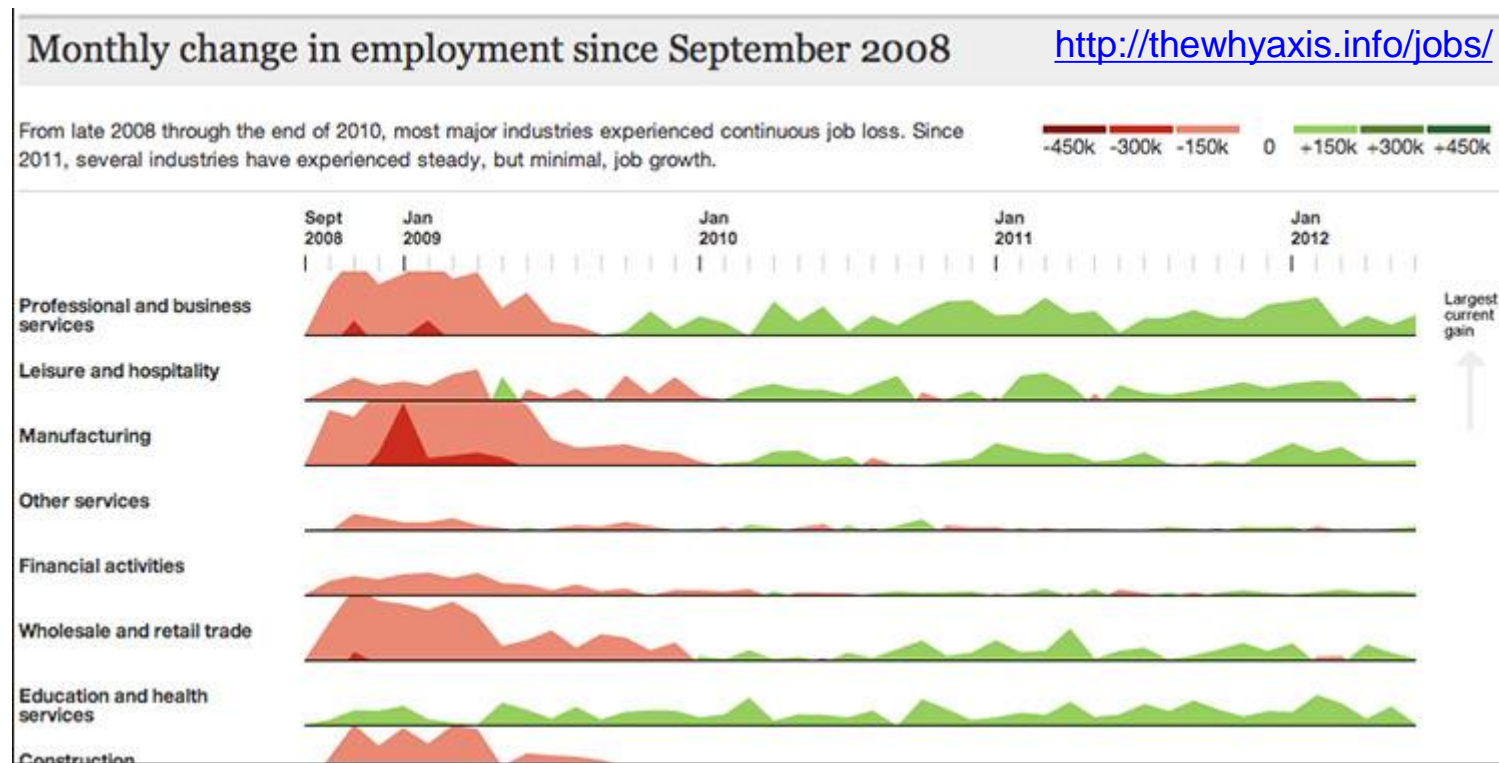
A time series is defined by the values  $y_1, y_2, \dots, y_n$  of variable  $Y$  at times  $t_1, t_2, \dots, t_n$ , thus  $y$  is a function of  $t$ :  $y = f(t)$ .

- ▶ **Measurement of water level in Nile River**



► **Time series analysis** aims to draw inferences, forecast future values, and control the series.

- Identifying the patterns of trends and seasonal variation in correlated data
- Understanding and modeling the data for explanation
- Predicting short-term trends based on previous patterns



## ► What are components of time series data?

### 1. Trend ( $T_t$ )

- A trend exists when there is a long-term increase or decrease in the data.

### 2. Seasonal variation ( $S_t$ )


- A seasonal pattern exists when a series is influenced by seasonal factors (ex: quarterly, monthly, half-yearly).

### 3. Cyclical variation ( $C_t$ )

- A cyclic pattern exists when data exhibit rises and falls that are not of fixed period.

### 4. Irregular or random variation ( $I_t$ )

- The variation of observations in a time series which is unusual or unexpected. Ex: Floods, fires, earthquakes, revolutions, epidemics, strikes, ...





## 2. Creating Time Series Objects

### (1) Creating a simple time series

The **ts()** function will convert a numeric vector into an R time series object

```
ts(data, start, end, frequency, ... ) {stats}
```


The function **ts** is used to create time-series objects.

frequency: the number of observations per unit time  
(1=annual, 4=quartly, 12=monthly, etc.)

```
> ## Using May 1954 as start date:  
> x <- 1:36; td <- runif(x)  
> mytsd <- ts(td, start = c(1954,7), frequency = 12)  
> mytsd
```

	Jan	Feb	Mar	Apr	May	Jun	Jul
1954						0.17064524	
1955	0.04766363	0.70085309	0.35188864	0.40894400	0.82095132	0.91885735	0.28252833
1956	0.47784538	0.56025326	0.69826159	0.91568354	0.61835123	0.42842151	0.54208037
1957	0.15288722	0.80341854	0.54682616	0.66231764	0.17169849	0.63305536	
	Aug	Sep	Oct	Nov	Dec		
1954	0.17217175	0.48204261	0.25296493	0.21625479	0.67437639		
1955	0.96110479	0.72839443	0.68637508	0.05284394	0.39522013		
1956	0.05847849	0.26085686	0.39715195	0.19774474	0.83192756		
1957							

```
> plot(mytsd) # using 'plot.ts' for time-series plot
```

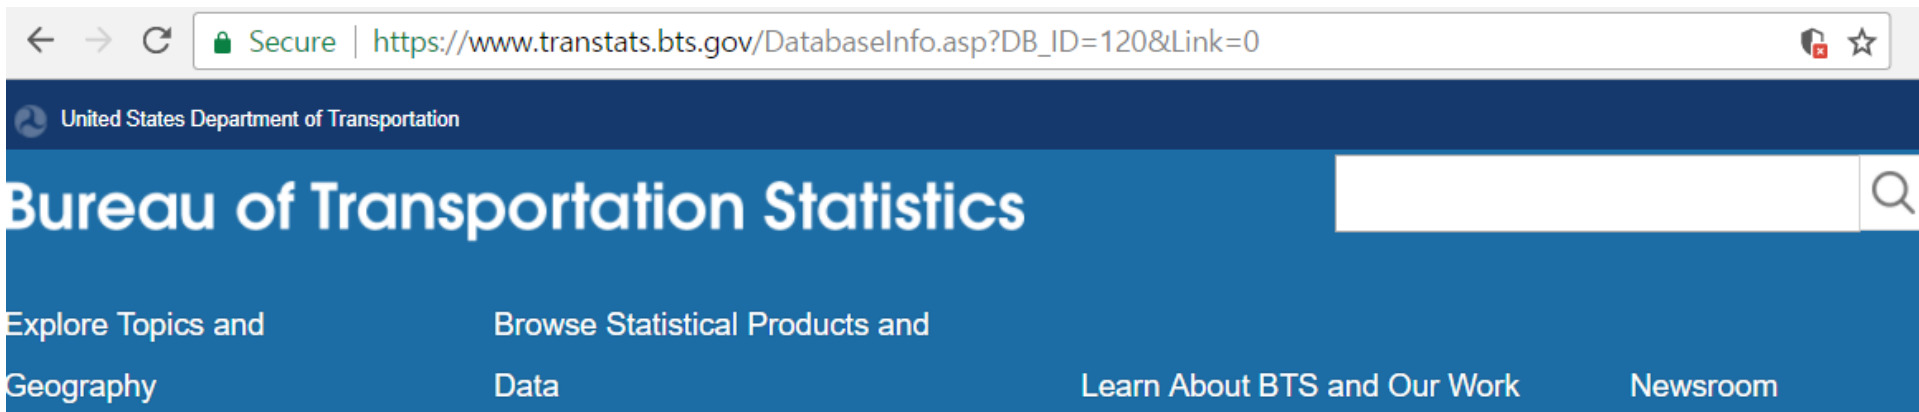


## (2) Sample Dataset: hflights

## 2. Creating Time Series Objects

**flights** {flights} Houston flights data

This dataset contains all flights departing from Houston airports IAH and HOU. Data from: [http://www.transtats.bts.gov/DatabseInfo.asp?DB\\_ID=120&Link=0](http://www.transtats.bts.gov/DatabseInfo.asp?DB_ID=120&Link=0)



```
> #(1) Reading time series data
> library(hflights)
> head(hflights,2)
```

	Year	Month	DayofMonth	DayOfWeek	DepTime	ArrTime	UniqueCarrier			
5424	2011	1	1	6	1400	1500	AA			
5425	2011	1	2	7	1401	1501	AA			
	FlightNum	TailNum	Actual	ElapsedTime	AirTime	ArrDelay	DepDelay	Origin		
5424	428	N576AA		60	40	-10	0	IAH		
5425	428	N557AA		60	45	-9	1	IAH		
	Dest	Distance	TaxiIn	TaxiOut	cancelled	cancellationCode	Diverted			
5424	DFW	224	7	13	0		0			
5425	DFW	224	6	9	0		0			6

Let's transform the data.frame to data.table for easy aggregation.

Then we create a **date** variable from the provided Year, Month, and DayofMonth columns.

**ISOdate**(year, month, day, ... ) {base}

Date-time conversion function from numeric representations

```
> library(data.table)
> dt <- data.table(hflights)
> dt[, date := ISOdate(Year, Month, DayofMonth)]
> head(dt, 2)
```

	Year	Month	DayofMonth	DayOfWeek	DepTime	ArrTime	UniqueCarrier	FlightNum	TailNum
1:	2011	1	1	6	1400	1500	AA	428	N576AA
2:	2011	1	2	7	1401	1501	AA	428	N557AA

	ActualElapsedTime	AirTime	ArrDelay	DepDelay	Origin	Dest	Distance	TaxiIn	TaxiOut
1:	60	40	-10	0	IAH	DFW	224	7	13
2:	60	45	-9	1	IAH	DFW	224	6	9

	Cancelled	CancellationCode	Diverted	date
1:	0		0	2011-01-01 12:00:00
2:	0		0	2011-01-02 12:00:00

date column was created!

We can compute **the total number of flights, the overall sum of arrival delays, the number of cancelled flights, and the average distance of the related flights** for each day in 2011.

```
> daily <- dt[, list( N = .N,
+   Delays      = sum(ArrDelay, na.rm=TRUE),
+   Cancelled   = sum(Cancelled),
+   Distance    = mean(Distance) ), by=date]
> head(daily)
```

Group by 'date' and count instances (.N)

	date	N	Delays	Cancelled	Distance
1:	2011-01-01 12:00:00	552	5507	4	827.0761
2:	2011-01-02 12:00:00	678	7010	11	786.7788
3:	2011-01-03 12:00:00	702	4221	2	772.3276
4:	2011-01-04 12:00:00	583	4631	2	754.5523
5:	2011-01-05 12:00:00	590	2441	3	759.5441
6:	2011-01-06 12:00:00	660	3994	0	755.5955

Ref: G. Daroczi, [Mastering Data Analysis with R](#) (Packt Pub., Birmingham, 2015) Chap. 12.



### 3. Time Series Visualization

- A time series plot is a graph that you can use to evaluate patterns and behavior in data over time.
- A time series plot displays observations on the y-axis against equally spaced time intervals on the x-axis.

```
setorder(x, ... ) {data.table}
```


Fast row reordering of a data.table by reference

```
ts(data, start, end, frequency, deltat, ... ) {stats}
```


The function ts is used to create time-series objects.

```
#sorting daily data by date  
setorder(daily,date)  
plot(ts(daily))
```

The x axis is indexed from 1 to 365 and the date transformed to timestamps on the y axis.



ts(daily)



# Decomposition of Time Series Data

Additive Model:  $Y_t = T_t + C_t + S_t + \epsilon_t$

Multiplicative Model:  $Y_t = T_t \times C_t \times S_t \times \epsilon_t$


```
decompose(x, type=c("additive", "multiplicative"), filter=NULL) {stats}
```

Decompose a time series into seasonal, trend and irregular components using moving averages. Deals with additive or multiplicative seasonal component.




**type = "additive"**

Time series = Seasonal + Trend + Random

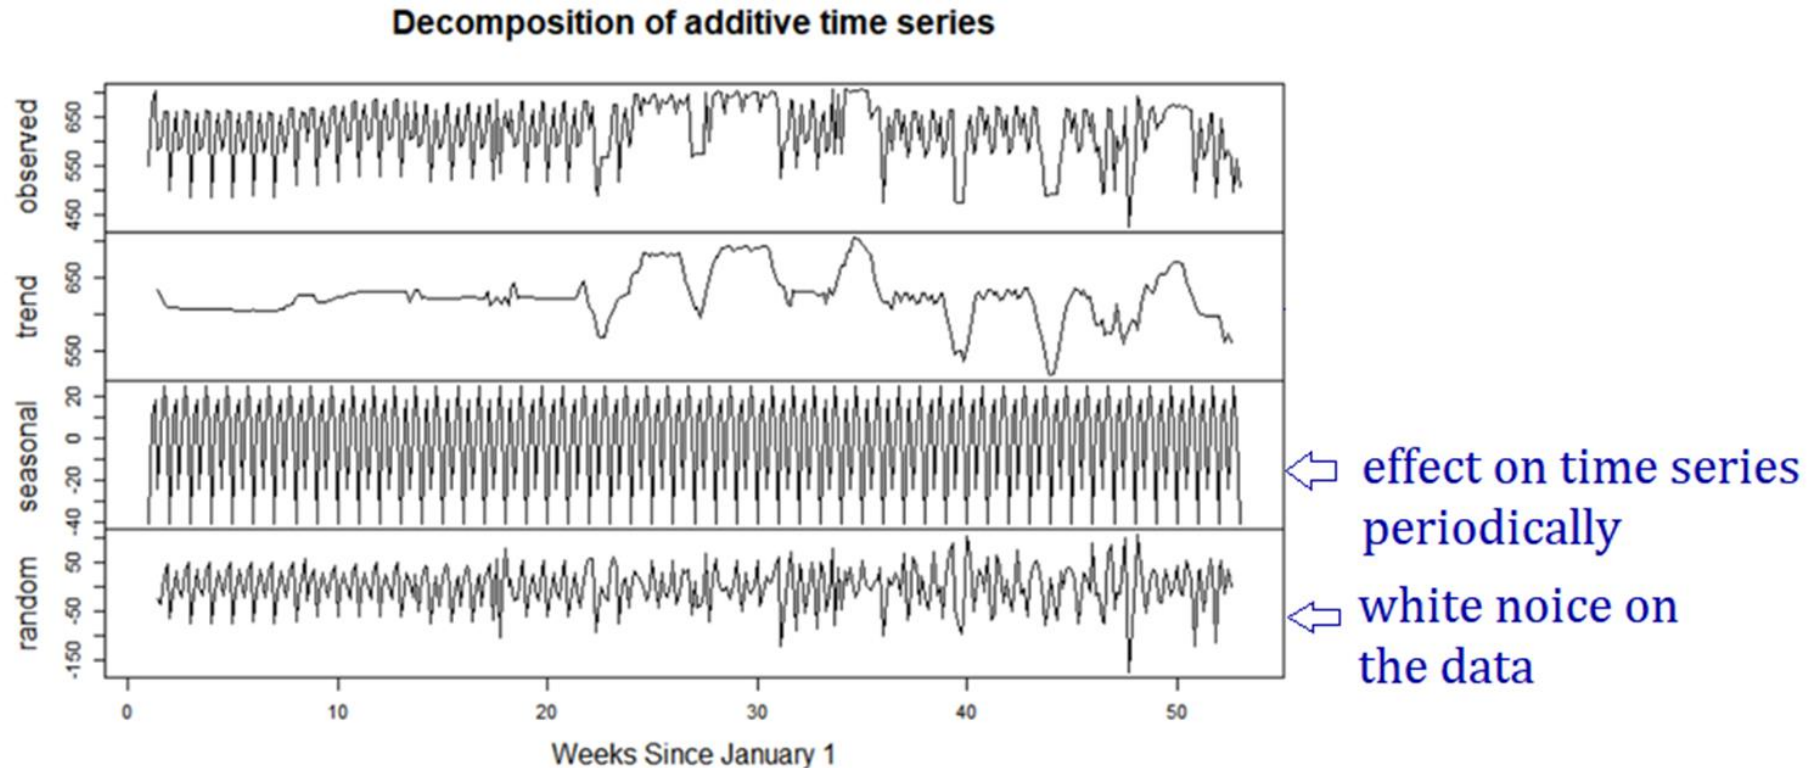


**type = "multiplicative"**

Time series = Seasonal \* Trend \* Random



```
daily7 <- decompose(ts(daily$N, frequency=7))  
plot(daily7)
```



In the **trend** plot: Peak interval between 25 and 35 weeks (summertime)  
Lowest number of flights on 46<sup>th</sup> week (probably due to Thanksgiving Day)

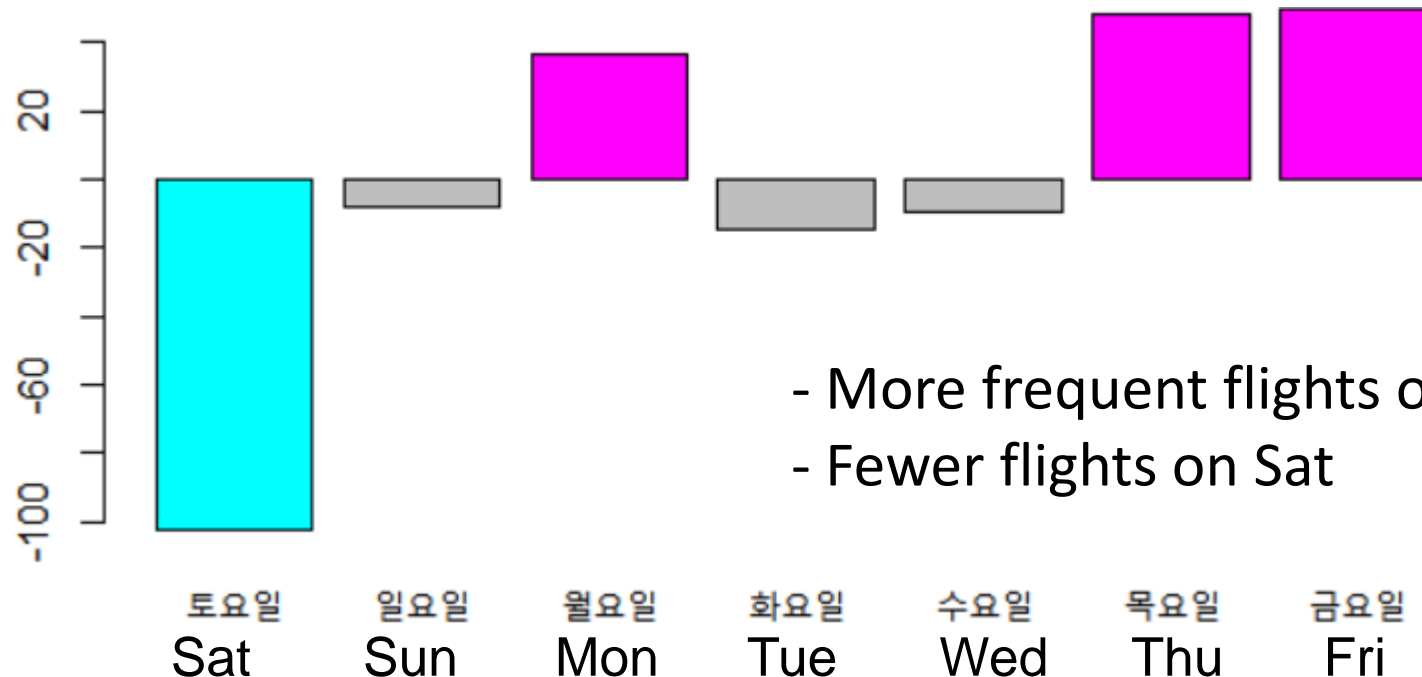
```
> weeks = setNames(daily7$figure, weekdays(daily$date[1:7]))
> weeks
```

Saturday	Sunday	Monday	Tuesday	Wednesday
토요일	일요일	월요일	화요일	수요일
-102.171776	-8.051328	36.595731	-14.928941	-9.483886

Thursday	Friday
목요일	금요일
48.335226	49.704974

```
barplot(weeks,col=c('cyan','gray','magenta',
                    'gray','gray','magenta','magenta'))
```



- More frequent flights on Mon, Thu, Fri
- Fewer flights on Sat

## 4. Time Series Stationarity

- **Stationarity** is prerequisite for time series modeling.
- Data in stationarity time series do not depend on time.
- So it does not have a trend or seasonality.

Three conditions of time series stationarity:

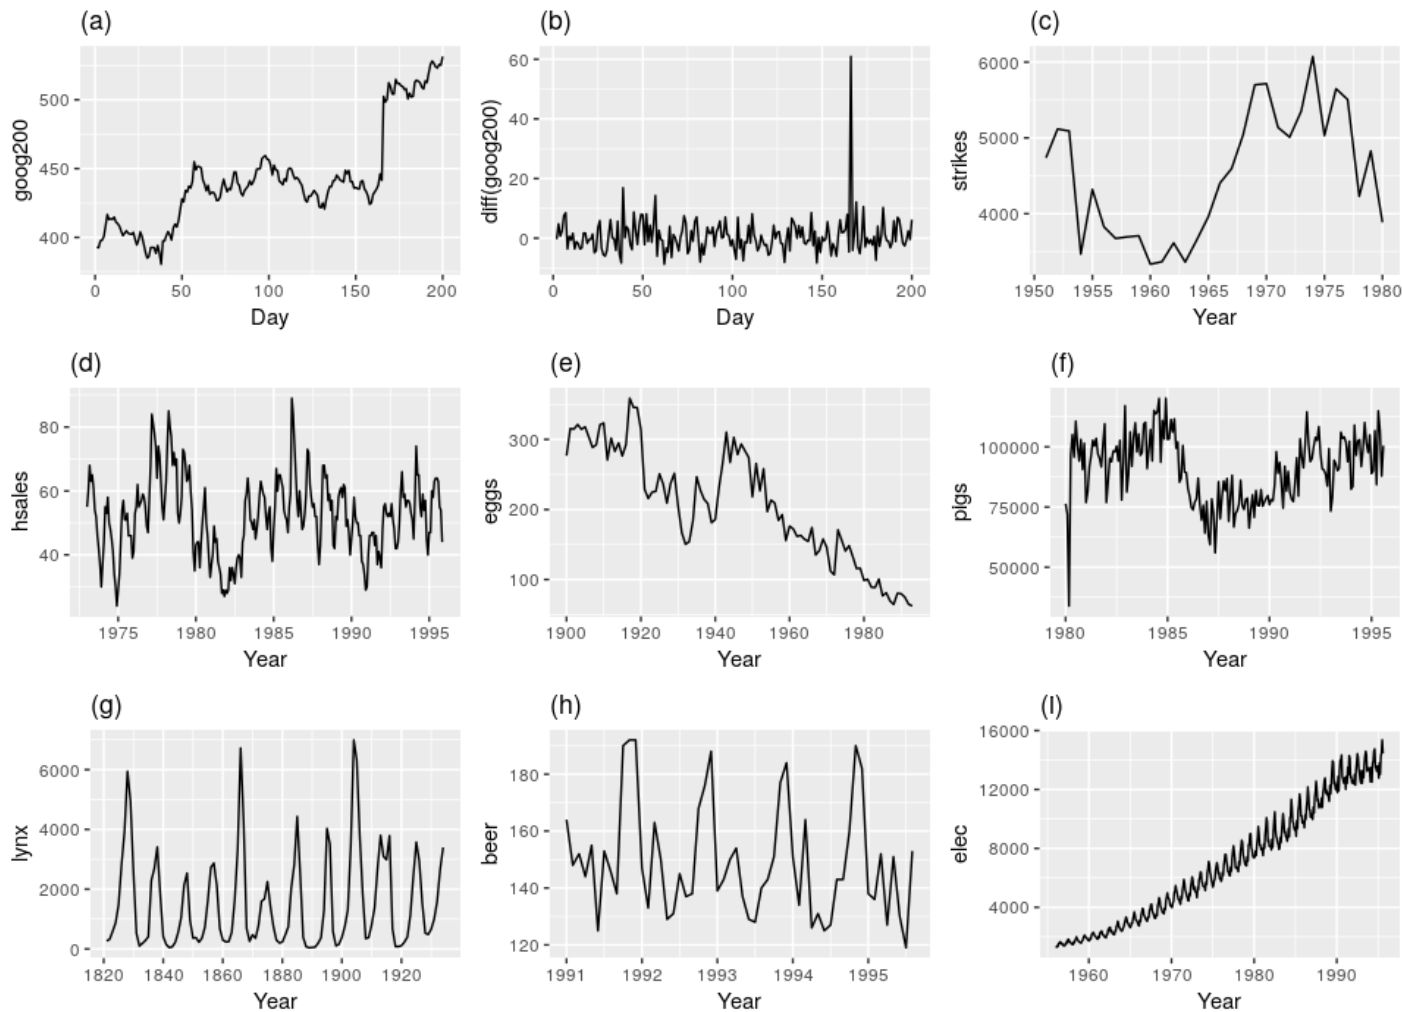
- (1) The **mean** value of time-series is constant over time, which implies, the trend component is nullified.
- (2) The **variance** does not increase over time.
- (3) **Seasonality** effect is minimal.

[Reference]

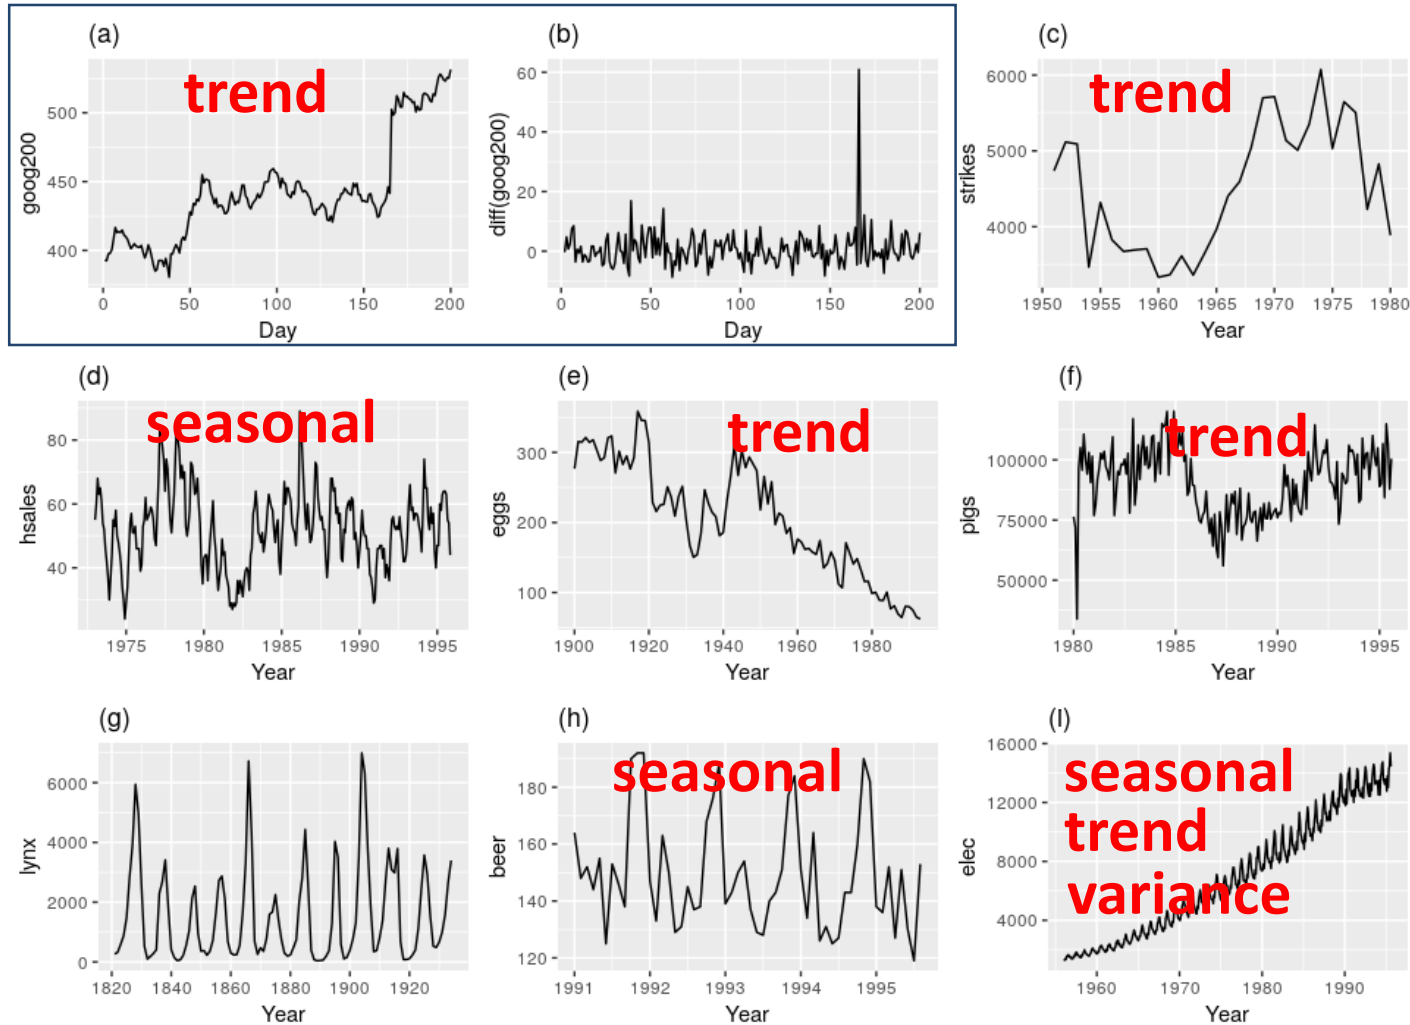
<http://r-statistics.co/Time-Series-Analysis-With-R.html>




## Which of these series are stationary?



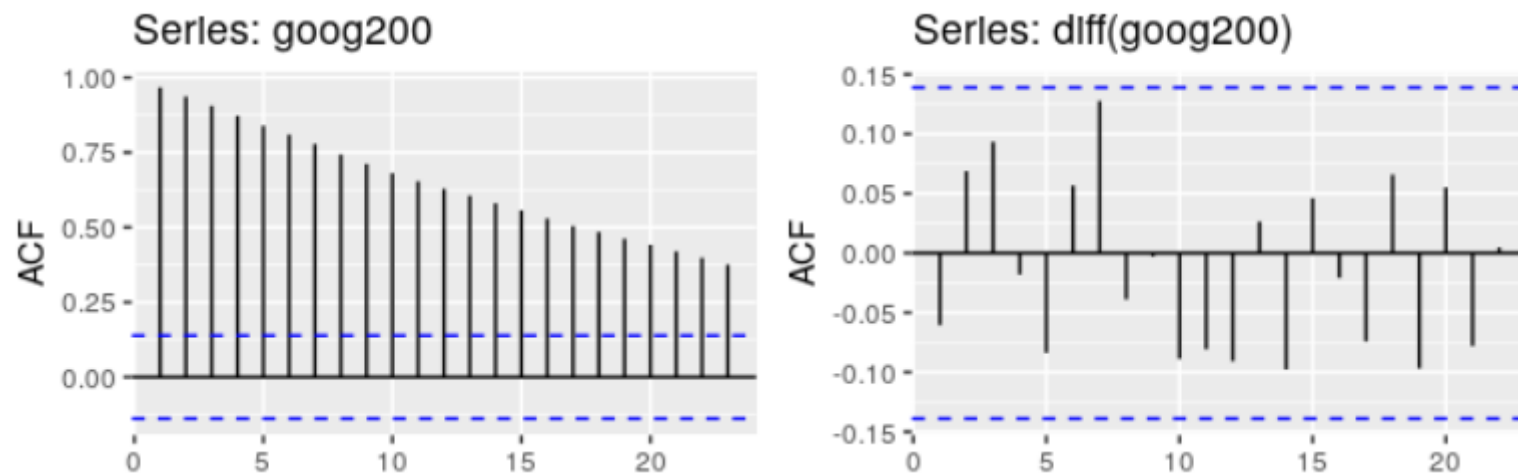
## Which of these series are stationary?



## Which of these series are stationary?



The raw stock price data in (a) is non-stationary, yet its daily price difference shown in (b) is stationary. This process is called **differencing**.




**ACF (autocorrelation function) plots** can be used to visualize this.

## AirPassengers {datasets}


Monthly Airline Passenger Numbers 1949-1960

The classic Box & Jenkins airline data. Monthly totals of international airline passengers, 1949 to 1960.

```
> class(AirPassengers)
[1] "ts"
> start(AirPassengers)
[1] 1949    1
> end(AirPassengers)
[1] 1960   12
> sm <- summary(AirPassengers); sm
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 104.0   180.0   265.5   280.3   360.5   622.0
> boxplot(sm, col='cyan')
> abline(h=sm[4])
```



```
plot(AirPassengers)
abline(reg=lm(AirPassengers~time(AirPassengers)),
       col='red')
```



Plotting the data shows: (1) the numbers are increasing and  
(2) there exists a seasonality pattern that repeats every year.



- **Autocorrelation Function (ACF)** plot shows the *correlation* of the series with *itself* at different lags. Autocorrelation of  $Y$  at lag  $h$  is the correlation between  $Y$  and  $LAG(Y, h)$

$$\gamma_X(t+h, t) = \text{Cov}(X_{t+h}, X_t)$$

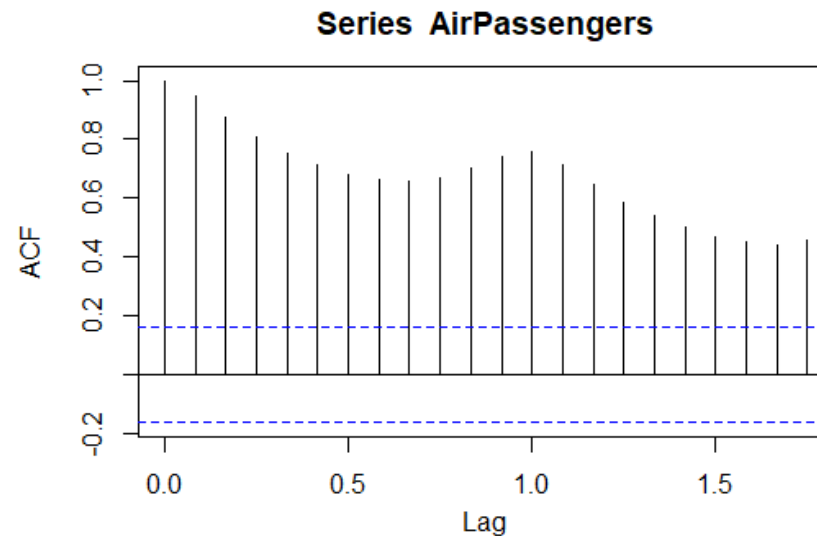
$$\rho_X(h) = \frac{\gamma_X(h)}{\gamma_X(0)} = \text{Corr}(X_{t+h}, X_t)$$

- **Partial autocorrelation function (PACF)** – measures correlation between (time series) observations that are  $k$  time periods apart after controlling for correlations at intermediate lags (i.e., lags less than  $k$ ). In other words, it is the correlation between  $Y_t$  and  $Y_{t-k}$  after removing the effects of intermediate  $Y$ 's.

# Autocorrelation Function

Autocorrelation function (ACF) is the correlation between series values that are  $k$  intervals apart at lag  $k$ .

```
# autocorrelation  
acf(AirPassengers)
```



In R `acf` starts with lag 0, that is the correlation of a value with itself. The **blue dashed lines** represent the confidence limits (insignificant zone)

## Identifying non-stationary series

- The ACF of stationary data drops to zero relatively quickly.
- The ACF of non-stationary data decreases slowly.

```
kpss.test(x, null = c("Level", "Trend")) {tseries}
```

KPSS Test for Stationarity

Computes the Kwiatkowski-Phillips-Schmidt-Shin (KPSS) test for the null hypothesis that  $x$  is level or trend stationary.

```
> kpss.test(AirPassengers)
```

```
      KPSS Test for Level Stationarity
```

```
data:  AirPassengers
```

```
KPSS Level = 2.7395, Truncation lag parameter = 4, p-value = 0.01
```

Interpretation: AirPassengers dataset is *not stationary* because it has both a trend and seasonality. The p-value for the test is less than 0.05. Thus we reject the null hypothesis of stationary.

# Differencing

- Differencing can help stabilize the mean of a time series by removing changes in the level of a time series, and therefore reducing trend and seasonality.
- The `diff()` function can make a non-stationary time series stationary.

**`diff(x, ...)`** {base}


Returns suitably lagged and iterated differences.

```
> ## Differencing
> AirPass_Diff <- diff(AirPassengers)
> kpss.test(AirPass_Diff)
```

KPSS Test for Level Stationarity

```
data: AirPass_Diff
KPSS Level = 0.014626, Truncation lag parameter = 4,
p-value = 0.1
```

```
plot(AirPass_Diff)
abline(h=0,col=2)
```



This result shows stationary.

- p-value = 0.1
- The mean and std variations have small variations with time.

## 5. Time Series and Forecasting

### (1) Reading Time Series Data

**co2** {datasets}

Mauna Loa Atmospheric CO2 Concentration

A time series of 468 observations; monthly from 1959 to 1997.

Atmospheric concentrations of CO2 are expressed in parts per million (ppm).


Source: C. D. Keeling and T. P. Whorf, Scripps Institution of Oceanography (SIO), University of California, La Jolla, California USA 92093-0220.

```
> co2 #time series data
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1959	315.42	316.31	316.50	317.56	318.13	318.00	316.39	314.65	313.68	313.18	314.66	315.43
1960	316.27	316.81	317.42	318.87	319.87	319.43	318.01	315.74	314.00	313.68	314.84	316.03
1961	316.73	317.54	318.38	319.31	320.42	319.61	318.42	316.63	314.83	315.16	315.94	316.85
1962	317.78	318.40	319.53	320.42	320.85	320.45	319.45	317.25	316.11	315.27	316.53	317.53
1963	318.58	318.92	319.70	321.22	322.08	321.31	319.58	317.61	316.05	315.83	316.91	318.20
1964	319.41	320.07	320.74	321.40	322.06	321.73	320.27	318.54	316.54	316.71	317.53	318.55
...												
1996	362.09	363.29	364.06	364.76	365.45	365.01	363.70	361.54	359.51	359.65	360.80	362.38
1997	363.23	364.06	364.61	366.40	366.84	365.68	364.52	362.57	360.24	360.83	362.49	364.34



```
# plot co2 time series  
plot(co2, xlab='Year', ylab='CO2 Concetnrnration (ppm)')
```




## (2) Seasonal Decomposition

```
stl(x, s.window, ... ) {stats}
```

Seasonal Decomposition of Time Series by Loess

Decompose a time series into seasonal, trend and irregular components using loess (locally weighted smoothing).

```
fit <- stl(co2, s.window="period")  
plot(fit)
```




### (3) Forecasting Time Series

We can now forecast future values based on the smoothing model.

```
library(forecast)
m <- Holtwinters(co2)
p <- predict(m, 50, prediction.interval = TRUE)
plot(m, p)
```

#### Holt-Winters filtering



The [getSymbols](#) function from the quantmod package is an easy and convenient way to bring historical stock prices into your R environment.

```
getSymbols(Symbols, src, auto.assign, ... ) {quantmod}  
Load and Manage Data from Multiple Sources
```

```
## TimeSeriesAnalysis5_financial.R ##  
library(quantmod)  
#Microsoft Corporation (Nasdaq, USA)  
Microsoft <- getSymbols("MSFT", src="yahoo", auto.assign=F,  
                        from='2015-10-01', to='2017-09-30')  
tail(Microsoft,2)
```

	MSFT.Open	MSFT.High	MSFT.Low	MSFT.Close	MSFT.Volume	MSFT.Adjusted
2017-09-28	73.54	73.97	73.31	73.87	10883800	72.84149
2017-09-29	73.94	74.54	73.88	74.49	17079100	73.45286

```
#Facebook, Inc. (Nasdaq, USA)  
Facebook <- getSymbols("FB", src="yahoo", auto.assign=F,  
                      from='2015-10-01', to='2017-09-30')  
tail(Facebook,2)
```

	FB.Open	FB.High	FB.Low	FB.Close	FB.Volume	FB.Adjusted
2017-09-28	167.94	169.07	167.16	168.73	12178700	168.73
2017-09-29	168.83	171.66	168.81	170.87	15340400	170.87

```

> #SKhynix (KOSPI, Republic of Korea)
> SKhynix <- getSymbols("000660.KS", src="yahoo", auto.assign=F,
+                       from='2017-12-11', to='2018-12-10')
> class(SKhynix)
[1] "xts" "zoo"
> tail(SKhynix)
      000660.KS.Open 000660.KS.High 000660.KS.Low
2018-12-03          71000          71100          69800
2018-12-04          69900          70400          68400
2018-12-05          67200          68800          67200
2018-12-06          67800          68200          65700
2018-12-07          66700          67400          66600
2018-12-10          65900          65900          64700
      000660.KS.Close 000660.KS.Volume 000660.KS.Adjusted
2018-12-03          70500          2760290          70500
2018-12-04          69000          3222060          69000
2018-12-05          68200          3008950          68200
2018-12-06          66000          3902275          66000
2018-12-07          66800          2590733          66800
2018-12-10          65500          2206084          65500

```




## (2) Time Series Plot

A time series plot is a graph that you can use to evaluate patterns and behavior in data over time.

A time series plot displays observations on the y-axis against equally spaced time intervals on the x-axis.

```
MSFT.Close = Microsoft[,4]  
plot(MSFT.Close) #close price
```



**chartSeries**(x, subset, theme, up.col, ... ) {quantmod}


Charting tool to create standard financial charts given a time series like object. Possible chart styles include candles, matches, bars, and lines.

```
#4=close 5=volume  
chartSeries(Microsoft[,4:5],subset="2015-10-01::2017-09-30",  
            theme='white', up.col='black')
```



If you already have the download Microsoft and Facebook, then you can merge the resulting closing prices.

```
#Comparing two stock prices  
z <- merge(Microsoft[,4], Facebook[,4])  
library("ggplot2")  
autoplot(z, facets = NULL)
```



### (3) Forecasting Time Series

```
auto.arima(y, test, ic, ... ) {forecast}
```

Fit best ARIMA model to univariate time series.

```
library(forecast)  
fit = auto.arima(Microsoft[,4], test='adf', ic='bic')  
plot(forecast(fit, h=48), main="Forecast by ARIMA of Microsoft")
```

**Forecast by ARIMA of Microsoft**

